

Homework_v1.r

nhirata

2020-01-29

```
# Question 5.1
# Test to see whether there are any outliers in the last column (number of crimes per 100,000 people).

# Dr. Sokol's advice:
# Check to see if the outlier is "Real" or an "Error".
# If Real, you need to research/investigate to see what caused it and decide whether your model needs to consider it or not.

rm(list = ls()) # Clear the list
library(outliers)
library(ggplot2)
#Load kknn and kernlab Libraries
set.seed(25)

#Read my credit card data
df <- read.table("C:/Users/nhirata/Desktop/Georgia Tech/OneDrive - Georgia Institute of Technology/Georgia Tech/ISYE_6501/Week_3/data 5.1/uscrime.txt", header=TRUE, stringsAsFactors = FALSE)

#Look at the first rows of my data to see if the data comes out right based on the parameters set in read.table.
head(df)
```

```
##      M So   Ed Po1 Po2   LF   M.F Pop   NW   U1 U2 Wealth Ineq   Prob
## 1 15.1  1  9.1  5.8  5.6 0.510 95.0  33 30.1 0.108 4.1   3940 26.1 0.084602
## 2 14.3  0 11.3 10.3  9.5 0.583 101.2  13 10.2 0.096 3.6   5570 19.4 0.029599
## 3 14.2  1  8.9  4.5  4.4 0.533 96.9  18 21.9 0.094 3.3   3180 25.0 0.083401
## 4 13.6  0 12.1 14.9 14.1 0.577 99.4 157  8.0 0.102 3.9   6730 16.7 0.015801
## 5 14.1  0 12.1 10.9 10.1 0.591 98.5  18  3.0 0.091 2.0   5780 17.4 0.041399
## 6 12.1  0 11.0 11.8 11.5 0.547 96.4  25  4.4 0.084 2.9   6890 12.6 0.034201
##      Time Crime
## 1 26.2011    791
## 2 25.2999   1635
## 3 24.3006    578
## 4 29.9012   1969
## 5 21.2998   1234
## 6 20.9995    682
```

```
summary(df)
```


```
##           M           So           Ed           Po1
## Min.      :11.90    Min.      :0.0000    Min.      : 8.70    Min.      : 4.50
## 1st Qu.:13.00    1st Qu.:0.0000    1st Qu.: 9.75    1st Qu.: 6.25
## Median :13.60    Median :0.0000    Median :10.80    Median : 7.80
## Mean      :13.86    Mean      :0.3404    Mean      :10.56    Mean      : 8.50
## 3rd Qu.:14.60    3rd Qu.:1.0000    3rd Qu.:11.45    3rd Qu.:10.45
## Max.      :17.70    Max.      :1.0000    Max.      :12.20    Max.      :16.60
##           Po2           LF           M.F           Pop
## Min.      : 4.100    Min.      :0.4800    Min.      : 93.40    Min.      : 3.00
## 1st Qu.: 5.850    1st Qu.:0.5305    1st Qu.: 96.45    1st Qu.: 10.00
## Median : 7.300    Median :0.5600    Median : 97.70    Median : 25.00
## Mean      : 8.023    Mean      :0.5612    Mean      : 98.30    Mean      : 36.62
## 3rd Qu.: 9.700    3rd Qu.:0.5930    3rd Qu.: 99.20    3rd Qu.: 41.50
## Max.      :15.700    Max.      :0.6410    Max.      :107.10    Max.      :168.00
##           NW           U1           U2           Wealth
## Min.      : 0.20    Min.      :0.07000    Min.      :2.000    Min.      :2880
## 1st Qu.: 2.40    1st Qu.:0.08050    1st Qu.:2.750    1st Qu.:4595
## Median : 7.60    Median :0.09200    Median :3.400    Median :5370
## Mean      :10.11    Mean      :0.09547    Mean      :3.398    Mean      :5254
## 3rd Qu.:13.25    3rd Qu.:0.10400    3rd Qu.:3.850    3rd Qu.:5915
## Max.      :42.30    Max.      :0.14200    Max.      :5.800    Max.      :6890
##           Ineq         Prob         Time         Crime
## Min.      :12.60    Min.      :0.00690    Min.      :12.20    Min.      : 342.0
## 1st Qu.:16.55    1st Qu.:0.03270    1st Qu.:21.60    1st Qu.: 658.5
## Median :17.60    Median :0.04210    Median :25.80    Median : 831.0
## Mean      :19.40    Mean      :0.04709    Mean      :26.60    Mean      : 905.1
## 3rd Qu.:22.75    3rd Qu.:0.05445    3rd Qu.:30.45    3rd Qu.:1057.5
## Max.      :27.60    Max.      :0.11980    Max.      :44.00    Max.      :1993.0
```

Quantitative Review:


Check for normality of the crime data since this is an assumption of the Grubbs test by using some visuals.

Null hypothesis: Data is normally distributed


`plot(df$Crime)`




```
plot(rep(0,length(df$Crime)),df$Crime)
```



```
plot(df$Crime,df$Pop) #based on this 2 dimensional plot between population and crime, we can see that the smallest population has the highest crime rate.
```




```
boxplot(df$Crime)
```



```
sort(boxplot(df$Crime, plot=FALSE)$out) #based on the boxplot, we can see that there are 3 outliers. But are they really outliers?
```


```
## [1] 1674 1969 1993
```

```
qqnorm(df$Crime);qqline(df$Crime) #based on the Q-Q test, the high outliers are key factors of the data not being normally distributed. But we should run a Shapiro test to be completely sure.
```




```
hist(df$Crime) #supplemental visual to see how the crime is distributed.
```


Histogram of df\$Crime



```
d <- density(df$Crime)
plot(d, main = "Density Plot")
polygon(d, col="red", border="blue") #supplemental visual to see the density of crime.
```

Density Plot



Before running the Grubbs.Test, we should always run a Q-Q Plot and Shapiro analysis to check if the underlying data is normally distributed because Grubbs.Test automatically assumes a normal distribution.
`shapiro.test(df$Crime)`

```
##  
## Shapiro-Wilk normality test  
##  
## data: df$Crime  
## W = 0.91273, p-value = 0.001882
```

```
# The test rejects the hypothesis of normality when the p-value is less than or equal to 0.05. Failing the normality test allows us to state with 95% confidence the data does not fit the normal distribution.  
# However, let's run the Grubbs Test anyway to see an initial view of what the p-values are per type (11,10, and 10 (opposite)) or respectively (both sides, one side, the other side).
```

```
#Null hypothesis: Not both the min and max points are outliers, but one could be.  
grubbs.test(df$Crime,type=11)
```

```
##  
## Grubbs test for two opposite outliers  
##  
## data: df$Crime  
## G = 4.26877, U = 0.78103, p-value = 1  
## alternative hypothesis: 342 and 1993 are outliers
```

```
# #Null hypothesis: No outlier in one tail or the other.  
grubbs.test(df$Crime, type=10)
```

```
##  
## Grubbs test for one outlier  
##  
## data: df$Crime  
## G = 2.81287, U = 0.82426, p-value = 0.07887  
## alternative hypothesis: highest value 1993 is an outlier
```

```
grubbs.test(df$Crime, type =10, opposite = TRUE) #opposite means the other side of the tail; No outlier in the other side of the tail.
```

```
##  
## Grubbs test for one outlier  
##  
## data: df$Crime  
## G = 1.45589, U = 0.95292, p-value = 1  
## alternative hypothesis: lowest value 342 is an outlier
```

```
# The alternative hypothesis, which we will conclude if we reject the null hypothesis, is that at the very least that most extreme point is an outlier (statistically).
# the p-values indicate that there is no evidence whatsoever that any of the data are outliers because the p-value is 1 or greatly above .05. If the p-value is greater than the significance level (0.05), the decision is to fail to reject the null hypothesis (meaning there is no outlier).
```

```
#Even if we remove the extreme outliers from the initial box plot data set so that Shapiro's p-value is greater 0.05, the Grubbs Test still shows that there are not outliers in the data because each p-value is greater than the significance level (0.05).
```

```
df_remove_outliers <- df # replicate dataframe so that it doesn't mess up the original.
```

```
crime <- df_remove_outliers[,16]
```

```
outliers <- boxplot(crime, plot=FALSE)$out
print(outliers) #the outliers
```

```
## [1] 1969 1674 1993
```

```
df_remove_outliers[which(crime %in% outliers),] #the outlier rows
```

```
##      M So  Ed Po1 Po2  LF  M.F Pop  NW  U1  U2 Wealth Ineq  Prob
## 4  13.6  0 12.1 14.9 14.1 0.577 99.4 157  8.0 0.102 3.9  6730 16.7 0.015801
## 11 12.4  0 10.5 12.1 11.6 0.580 96.6 101 10.6 0.077 3.5  6570 17.0 0.016201
## 26 13.1  0 12.1 16.0 14.3 0.631 107.1  3  7.7 0.102 4.1  6740 15.2 0.041698
##      Time Crime
## 4  29.9012 1969
## 11 41.6000 1674
## 26 22.1005 1993
```

```
df_remove_outliers <- df_remove_outliers[-which(crime %in% outliers),] #remove the outlier rows
shapiro.test(crime) #run the shapiro test again
```

```
##
## Shapiro-Wilk normality test
##
## data:  crime
## W = 0.91273, p-value = 0.001882
```

```
#Null hypothesis: Not both the min and max points are outliers, but one could be.  
grubbs.test(crime,type=11)
```

```
##  
## Grubbs test for two opposite outliers  
##  
## data: crime  
## G = 4.26877, U = 0.78103, p-value = 1  
## alternative hypothesis: 342 and 1993 are outliers
```

```
# #Null hyposthesis: No outlier in one tail.  
grubbs.test(crime, type=10)
```


```
##  
## Grubbs test for one outlier  
##  
## data: crime  
## G = 2.81287, U = 0.82426, p-value = 0.07887  
## alternative hypothesis: highest value 1993 is an outlier
```

```
grubbs.test(crime, type =10, opposite = TRUE) #opposite means the other side of the tail; No outlier in the other side of the tail.
```

```
##  
## Grubbs test for one outlier  
##  
## data: crime  
## G = 1.45589, U = 0.95292, p-value = 1  
## alternative hypothesis: lowest value 342 is an outlier
```


```
#Another approach (Use a log scale)
#Based on the analysis provided by http://www.statsci.org/data/general/uscrime.html (where we retrieve the uscrime data), it
states that, "Crime is slightly better modeled on a log scale".
# So Lets take the log scale of crime and see what happens.
df["log_crime"]<-log(df$Crime) # New column for log scale crime.

#The analysis below shows that you can actually transform the data to make it a normal distribution.
boxplot(df$log_crime) #no outliers based on log scaled crime
```




```
qqnorm(df$log_crime);qqline(df$log_crime) #based on the Q-Q test, there are no outliers for the log scaled crime. But we sho
uld run a Shapiro test to be completely sure.
```

Normal Q-Q Plot




```
hist(df$log_crime) #supplemental visual to see how the log scaled crime is distributed.
```

Histogram of df\$log_crime



```
d <- density(df$log_crime)
plot(d, main = "Density Plot")
polygon(d, col="red", border="blue") #supplemental visual to see the density of log scaled crime.
```


Density Plot



```
# comparison between log scale and original data.  
set <-df[order(-df$log_crime, df$Crime),]  
head(set[,16:17],1)
```

```
##      Crime log_crime  
## 26  1993  7.597396
```

```
tail(set[,16:17],1)
```

```
## Crime log_crime
## 27 342 5.834811
```

```
#Null hypothesis: Not both the min and max points are outliers, but one could be.
grubbs.test(df$log_crime,type=11) # 7.597... refers to the highest crime rate and 5.834... refers to the lowest crime rate.
```

```
##
## Grubbs test for two opposite outliers
##
## data: df$log_crime
## G = 4.28791, U = 0.80013, p-value = 1
## alternative hypothesis: 5.8348107370626 and 7.59739632021279 are outliers
```

```
# #Null hyposthesis: No outlier in one tail.
grubbs.test(df$log_crime, type=10)
```

```
##
## Grubbs test for one outlier
##
## data: df$log_crime
## G = 2.16544, U = 0.89585, p-value = 0.6329
## alternative hypothesis: lowest value 5.8348107370626 is an outlier
```

```
grubbs.test(df$log_crime, type =10, opposite = TRUE) #opposite means the other side of the tail; No outlier in the other side of the tail.
```

```
##
## Grubbs test for one outlier
##
## data: df$log_crime
## G = 2.12247, U = 0.89994, p-value = 0.712
## alternative hypothesis: highest value 7.59739632021279 is an outlier
```

#Because each grubbs.test has a p-value greater than 0.05, we fail to reject the null hypothesis (meaning there is no outlier)

#Qualitative Review: (Deeper Investigation)

The smallest population state (300K) had the HIGHEST Crime Rate and the 2nd highest population state came in 2nd.

This seemed odd (especially when the lowest crime rate state had 2x more population than the highest crime rate state) so I researched crime statistics per state from <http://www.disastercenter.com/crime/> from years 1960 to 1965. Please see the attached "US. States Crime Rate per 100,000.xlsx" excel workbook for the analysis.

I discovered that Nevada and California had roughly the same population range and crime rate statistics compared to the top 2 crime states in our dataset.

Even though Nevada is one of the smallest population states, it's crime rate is consistently ranked #1 or #2 against California. I attached a U.S. Crime statistic report for year 1960-1965 w/ conditional formatting to easily see how Nevada compares to the other states.

Because Nevada and California's crime rate stays consistent each year, it should definitely be included in the data set.

Dr. Sokol taught us that there are outliers that are "real" (weird but consistent through out time) or an "error" (ex. hitting an extra "0" on your keyboard; mistake). I believe this is a "real" outlier

Sometimes statistical modeling can only go so far (not everything is black and white) thus thorough research/qualitative analysis must have a hand in providing a complete answer to certain questions.

Question 6.1

Describe a situation or problem from your job, everyday life, current events, etc., for which a Change Detection model would be appropriate. Applying the CUSUM technique, how would you choose the critical value and the threshold?

At the Jet Propulsion Laboratory (JPL), I work as a Business Analyst for the Mission Systems Engineering Section. JPL mainly uses change detection when measuring a satellite/rover's trajectory to make sure it's going the right way.

We have engineers in Mission Control that receive alerts when a certain project goes offcourse and the alerts they set are based on a change detection model. Based on CUSUM, I would initially choose a critical value based on 1/2 times the standard deviation of the mean and a T-value at 5 times the standard deviation of the mean. However, I would adjust the T-value again based on how sensitive the outcome. For this example, I would want a smaller target value than normal because I would not want to be responsible for crashing a multi-million dollar project (meaning, I'll take the false alarms).

Question 6.2.1 & 6.2.2

I tried to attempt answering in R but i wasn't able to figure out how to loop through each column efficiently (The code below only goes up to 1997).

Therefore, I created an excel workbook that was able to transparently model and present the data. Workbook title is "CUSUM 6.2"

I provided all my answers for Question 6.2.1 & 6.2.2 in the excel workbook but feel free to look below at my attempt in R.

Attempt at programming in R.(Used Excel Workbook instead)

```
rm(list = ls()) # Clear the list
# library(outliers)
# library(ggplot2)
# Load kkn and kernlab Libraries
set.seed(25)
temps <- read.table("C:/Users/nhirata/Desktop/Georgia Tech/OneDrive - Georgia Institute of Technology/Georgia Tech/ISYE_6501/Week_3/data 6.2/temps.txt", header=TRUE, stringsAsFactors = FALSE)


#####1996#####
# average the temperature for each column year
month_rows <- temps[1:31,]
month_avgs <- colMeans(month_rows[c(2:length(month_rows))]), dims=1, na.rm=T)
month_avgs
```

```
##      X1996      X1997      X1998      X1999      X2000      X2001      X2002      X2003
## 91.19355 87.25806 89.70968 87.64516 91.74194 86.74194 89.25806 85.58065
##      X2004      X2005      X2006      X2007      X2008      X2009      X2010      X2011
## 87.83871 86.93548 90.19355 86.41935 89.16129 86.64516 91.25806 91.93548
##      X2012      X2013      X2014      X2015
## 94.09677 84.70968 86.61290 90.06452
```

```
# compute the mean of the (now averaged) time series
da_mu <- mean(month_avgs)
# compute the difference between the mean of the time series and each "day"
da_minus_mu <- month_avgs[c(1)]-temps$X1996
# set C
C <- 1.95429 #1/2 times standard deviation (found this values in hindsight after excel)
t <- 19.5429 #5 times standard deviation (found this values in hindsight after excel)
# subtract C from the difference score
damimu_minus_C <- da_minus_mu - C
# create an empty vector for looping
# include an additional zero to help with indexing
cusum <- 0 * damimu_minus_C

# Loop through each day, check the cumulative sum, update the
# index of our accumulator with the appropriate value
#X1996
for (i in 1:length(damimu_minus_C))
{
  checker <- cusum[i] + damimu_minus_C[i+1]

  ifelse(checker > 0, cusum[i+1] <- checker, cusum[i+1] <- 0)
}
plot(cusum)
```



cusum

```
## [1] 0.0000000 0.0000000 0.0000000 0.0000000 0.2392584 0.0000000
## [7] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
## [13] 0.0000000 0.0000000 7.2392584 5.4785168 0.0000000 0.0000000
## [19] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
## [25] 5.2392584 10.4785168 17.7177752 27.9570335 27.1962919 25.4355503
## [31] 27.6748087 30.9140671 30.1533255 35.3925839 33.6318423 29.8711006
## [37] 31.1103590 29.3496174 34.5888758 33.8281342 34.0673926 35.3066510
## [43] 38.5459094 43.7851677 47.0244261 47.2636845 46.5029429 44.7422013
## [49] 42.9814597 42.2207181 42.4599765 41.6992348 39.9384932 38.1777516
## [55] 36.4170100 41.6562684 42.8955268 48.1347852 51.3740435 52.6133019
## [61] 57.8525603 65.0918187 74.3310771 90.5703355 92.8095939 98.0488523
## [67] 100.2881106 100.5273690 100.7666274 101.0058858 99.2451442 104.4844026
## [73] 107.7236610 108.9629194 120.2021777 130.4414361 133.6806945 140.9199529
## [79] 148.1592113 159.3984697 169.6377281 179.8769865 191.1162448 199.3555032
## [85] 204.5947616 209.8340200 212.0732784 217.3125368 227.5517952 241.7910535
## [91] 259.0303119 284.2695703 307.5088287 324.7480871 329.9873455 349.2266039
## [97] 372.4658623 397.7051206 426.9443790 438.1836374 457.4228958 474.6621542
## [103] 494.9014126 515.1406710 531.3799294 541.6191877 549.8584461 559.0977045
## [109] 566.3369629 589.5762213 615.8154797 637.0547381 647.2939965 655.5332548
## [115] 675.7725132 692.0117716 708.2510300 722.4902884 736.7295468 744.9688052
## [121] 752.2080635 759.4473219 767.6865803
```


which(cusum >19.5429) #the first number represents the row number of when the trigger activates based on the T value which would be July 28th.

```
## [1] 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46
## [20] 47 48 49 50 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65
## [39] 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80 81 82 83 84
## [58] 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 100 101 102 103
## [77] 104 105 106 107 108 109 110 111 112 113 114 115 116 117 118 119 120 121 122
## [96] 123
```

```
#####1997#####
month_avgs[c(2)]
```

```
## X1997
## 87.25806
```

```
da_minus_mu <- month_avgs[c(2)]-temps$X1997
damimu_minus_C <- da_minus_mu - C
cusum <- 0 * damimu_minus_C
for (i in 1:length(damimu_minus_C))
{
  checker <- cusum[i] + damimu_minus_C[i+1]
  ifelse(checker > 0, cusum[i+1] <- checker, cusum[i+1] <- 0)
}
plot(cusum)
```



cusum


```
## [1] 0.000000 0.000000 0.000000 0.000000 1.303775 2.607549
## [7] 12.911324 11.215098 12.518873 10.822647 12.126422 9.430196
## [13] 8.733971 4.037745 0.000000 0.000000 0.000000 0.000000
## [19] 0.000000 0.000000 0.000000 1.303775 0.000000 0.000000
## [25] 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000
## [31] 13.303775 18.607549 19.911324 17.215098 13.518873 10.822647
## [37] 12.126422 13.430196 18.733971 31.037745 36.341520 35.645294
## [43] 32.949069 30.252843 28.556618 25.860392 20.164167 14.467941
## [49] 10.771716 7.075490 4.379265 7.683039 13.986814 18.290588
## [55] 21.594363 22.898137 21.201912 16.505686 11.809461 6.113235
## [61] 0.417010 0.000000 0.000000 0.000000 0.000000 4.303775
## [67] 8.607549 11.911324 11.215098 8.518873 9.822647 15.126422
## [73] 18.430196 17.733971 16.037745 14.341520 11.645294 8.949069
## [79] 4.252843 1.556618 0.000000 0.000000 0.000000 15.303775
## [85] 20.607549 23.911324 43.215098 58.518873 79.822647 97.126422
## [91] 105.430196 104.733971 115.037745 127.341520 137.645294 144.949069
## [97] 149.252843 152.556618 155.860392 159.164167 164.467941 167.771716
## [103] 171.075490 177.379265 182.683039 199.986814 222.290588 250.594363
## [109] 269.898137 291.201912 307.505686 322.809461 338.113235 361.417010
## [115] 383.720785 407.024559 417.328334 431.632108 459.935883 490.239657
## [121] 511.543432 530.847206 556.150981
```

```
which(cusum >19.5429) # August 2nd
```

```
## [1] 33 40 41 42 43 44 45 46 47 55 56 57 85 86 87 88 89 90 91
## [20] 92 93 94 95 96 97 98 99 100 101 102 103 104 105 106 107 108 109 110
## [39] 111 112 113 114 115 116 117 118 119 120 121 122 123
```

```
#Thank you for taking the time to look through this code and my Homework assignment.
```