

ISYE 7406 Homework 5

October 24, 2021

Introduction

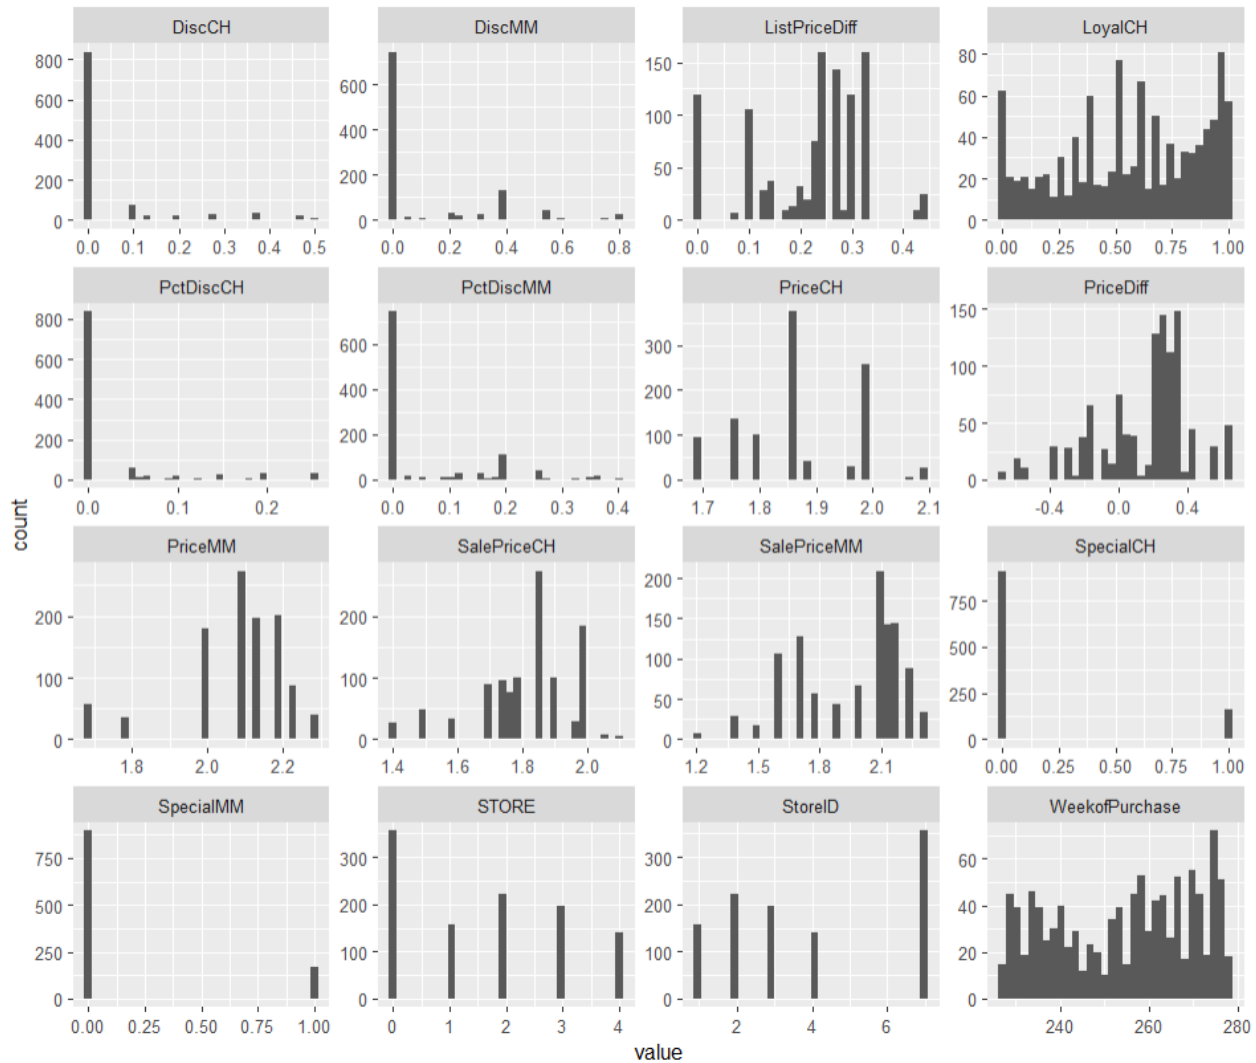
This report will review the orange juice (OJ) dataset which includes the *Purchase* response (y) variable and predictor (x) variables (WeekofPurchase, StoreID, PriceCH, PriceMM, DiscCH, DiscMM) and other variables with their descriptions that can be found [here](#). The purpose of analyzing the dataset is to fit a classification tree with “gini” criterion to the training and test sets to identify the Training & Test errors while also evaluating the results in a confusion matrix. Then we determine the optimal tree size that corresponds to the lowest cross-validation classification error rate and produce a few pruned trees corresponding to the optimal tree size obtained using cross validation and evaluated their performance on the test error and confusion matrix. The purpose of this report is to also identify other findings and visually explore the dataset for important/unusual patterns.

Exploratory Data Analysis

The exploratory data analysis will check for important/unusual patterns in the data through plots and summary statistics. According to the summary statistics shown below, it was interesting to see how PriceMM (price of Minute Maid orange juice) is generally pricier than PriceCH (Citrus Hill orange juice) suggesting that most of the purchases (~64%) are Citrus Hill (CH). It's also worth noting that StoreID, SpecialCH, SpecialMM, and STORE are categorical variables so I made sure to change the types to “factors” shown below but I changed them back to numeric for modeling purposes. Note: The 0 and 1's in Purchase represent CH and MM respectively.

Purchase	WeekofPurchase	StoreID	PriceCH	PriceMM	DiscCH
0:653	Min. :227.0	Min. :1.00	Min. :1.690	Min. :1.690	Min. :0.00000
1:417	1st Qu.:240.0	1st Qu.:2.00	1st Qu.:1.790	1st Qu.:1.990	1st Qu.:0.00000
	Median :257.0	Median :3.00	Median :1.860	Median :2.090	Median :0.00000
	Mean :254.4	Mean :3.96	Mean :1.867	Mean :2.085	Mean :0.05186
	3rd Qu.:268.0	3rd Qu.:7.00	3rd Qu.:1.990	3rd Qu.:2.180	3rd Qu.:0.00000
	Max. :278.0	Max. :7.00	Max. :2.090	Max. :2.290	Max. :0.50000
DiscMM	SpecialCH	SpecialMM	LoyalCH	SalePriceMM	
Min. :0.0000	Min. :0.0000	Min. :0.0000	Min. :0.000011	Min. :1.190	
1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.325257	1st Qu.:1.690	
Median :0.0000	Median :0.0000	Median :0.0000	Median :0.600000	Median :2.090	
Mean :0.1234	Mean :0.1477	Mean :0.1617	Mean :0.565782	Mean :1.962	
3rd Qu.:0.2300	3rd Qu.:0.0000	3rd Qu.:0.0000	3rd Qu.:0.850873	3rd Qu.:2.130	
Max. :0.8000	Max. :1.0000	Max. :1.0000	Max. :0.999947	Max. :2.290	
SalePriceCH	PriceDiff	Store7	PctDiscMM	PctDiscCH	ListPriceDiff
Min. :1.390	Min. : -0.6700	0:714	Min. :0.0000	Min. :0.00000	Min. :0.000
1st Qu.:1.750	1st Qu.: 0.0000	1:356	1st Qu.:0.0000	1st Qu.:0.00000	1st Qu.:0.140
Median :1.860	Median : 0.2300		Median :0.0000	Median :0.00000	Median :0.240
Mean :1.816	Mean : 0.1465		Mean :0.0593	Mean :0.02731	Mean :0.218
3rd Qu.:1.890	3rd Qu.: 0.3200		3rd Qu.:0.1127	3rd Qu.:0.00000	3rd Qu.:0.300
Max. :2.090	Max. : 0.6400		Max. :0.4020	Max. :0.25269	Max. :0.440
STORE					
Min. :0.000					
1st Qu.:0.000					
Median :2.000					
Mean :1.631					
3rd Qu.:3.000					
Max. :4.000					

Based on the distributions of the data (shown below), we can see that the discounts are somewhat similar but we can see that Minute Maid (MM) provided higher ranges of discounts compared to Citrus Hill (CH). You can also see from the StoreID distribution that most of the sales are from Store7 which is probably why it has its own predictor variable (Store7). We can also see that the sale prices are generally higher for (MM) compared to (CH) as we previously stated but there were instances where (MM) is cheaper than (CH) when on sale. But again, as we see from the PriceDiff distribution (Sale price of MM less sale price of CH), the distribution is left tailed meaning that (MM) are more expensive.



[illegible]

Methodology – Modeling for Training and Test errors

I performed the classification decision tree w/ “gini” criterion to retrieve training and test errors when predicting Purchase (CH or MM). The testing data is used to evaluate the model’s true performance since the training data may have excluded/included certain data that may have a significant positive/negative affect on the training error. For this exercise, I performed a 100-run cross validation for the model to obtain the test errors and observed which one had the best test error. Running a Monte Carlo cross-validation helps reduce the variance compared to just running one test run and depicts which model is truly the most accurate on a large scale.


Results & Findings

Classification tree w/ gini criterion – The “Gini” criterion helps calculate the probability of attributes that are classified incorrectly when randomly selected with the degree of Gini index varying between 0 and 1. Based on the summary statistics, we can see *complexity table* (shown below) that provides information about all of the trees of the model. The table displays the complexity parameter (CP), number of splits (np), resolution error rate (rel error), cross-validated error rate (xerror) and associated standard error (xstd). The complexity parameter (CP) determines whether a split is needed based on increasing the model’s fit and it’s main role is to save computation time by pruning off splits that don’t benefit the model. The cross-validated error rate (xerror) shows which tree is optimal based on the lowest xerror which as shown below is tree #5 with xerror 0.5209003 with 7 splits. It’s also worth noting that the default cross validation is a 10-fold.

```
Call:
rpart(formula = Purchase ~ ., data = ojtrain, method = "class",
      parms = list(split = "gini"))
n= 800
```

	CP	nsplit	rel error	xerror	xstd
1	0.5000000	0	1.000000	1.00000	0.04330127
2	0.0218750	1	0.500000	0.51250	0.03568252
3	0.0187500	2	0.478125	0.52500	0.03600130
4	0.0140625	4	0.440625	0.50625	0.03551999
5	0.0125000	6	0.412500	0.47500	0.03467483
6	0.0100000	8	0.387500	0.47500	0.03467483

When running the classification tree (shown below) the 0 and 1's in the circles represent CH and MM respectively. It's not surprising to see that the first branches came from the LoyalCH variable since we knew that the attribute was the only one that had a high correlation with Purchase. We can also see that there are 9 terminal nodes on the tree (outlined in red). Terminal nodes contain small subsets of the data where splitting is no longer valuable. The tree below shows that most of the CH values can be classified by $LoyalCH \geq .5086$ and $PriceDiff \geq -0.39$ which accounted for a count of 375/480 (~78%) but MM as well as the remaining CH values had to be split several times to get to their terminal nodes. This ultimately led to a training error result at 0.155 which is quite high compared to past models we've created.




Confusion Matrix and Statistics		
Prediction	Reference	
	0	1
0	150	23
1	23	74
Accuracy : 0.8296		
95% CI : (0.7794, 0.8725)		
No Information Rate : 0.6407		
P-Value [Acc > NIR] : 5.945e-12		
Kappa : 0.6299		
McNemar's Test P-Value : 1		
Sensitivity : 0.7629		
Specificity : 0.8671		
Pos Pred Value : 0.7629		
Neg Pred Value : 0.8671		
Prevalence : 0.3593		
Detection Rate : 0.2741		
Detection Prevalence : 0.3593		
Balanced Accuracy : 0.8150		
'Positive' Class : 1		

We can also see based on the confusion matrix (shown on the left) how the classification model performed. Accuracy resulted at 82.96% with a Sensitivity of 76.29% and Specificity of 86.71%. A low Sensitivity means that there are more FALSE negative results meaning more cases of Purchase being misclassified as 1. Specificity being higher shows that the model tends give less FALSE positive results meaning less cases of the model stating 1 when the true value is 0. The test set resulted in a test error of 0.1703704 which was higher than the training set.

Going back to the summary output from the training set, we can see that the optimal tree size that corresponds to the lowest cross-validation is tree #5 with 6 splits based on the smallest xerror (shown below).

```
Call:
rpart(formula = Purchase ~ ., data = ojtrain, method = "class",
      parms = list(split = "gini"))
n= 800
```

	CP	nsplit	rel error	xerror	xstd
1	0.5000000	0	1.000000	1.00000	0.04330127
2	0.0218750	1	0.500000	0.51250	0.03568252
3	0.0187500	2	0.478125	0.52500	0.03600130
4	0.0140625	4	0.440625	0.50625	0.03551999
5	0.0125000	6	0.412500	0.47500	0.03467483
6	0.0100000	8	0.387500	0.47500	0.03467483



Therefore, it makes sense to prune the tree based on the corresponding optimal tree size obtained above to achieve a better test error. After computing the pruned tree that corresponds to the optimal tree size model with $cp = 0.0125$, the test error resulted with a lower test error of 0.1666667 and had only 7 terminal nodes. Just in case, I attempted to run models with each indicated cp point on the above graph to identify all the results (shown below) and it showed that $cp = 0.014063$ and 0.01875 had the best test error of 0.155556 which was a better score than not only the unpruned tree but also the pruned tree with $cp = 0.0125$. It was interesting to see how the best cp values were extremely close to the horizontal line drawn 1 SE above the minimum of the curve. This makes sense since the best choice is often the leftmost value that lies below the horizontal line. It's also interesting to see once the cp value is a little above the line, the test error performs a lot worse.

cp	0.01	0.0125	0.014063	0.01875	0.021875	0.5
Test Error	0.17037	0.166667	0.155556	0.155556	0.2	0.359259

Conclusion

When analyzing all the cross-validated xerrors, it is clear that the pruned trees with cp equaling 0.014063 and 0.01875 performed better than the unpruned & optimized cross-validated prune tree. However, we can also see that test errors for the pruned trees are slightly higher than past classification models we've ran so it's worth investigating a model like Logistic Regression to see how it compares. The test error for the Logistic Regression resulted at 0.1444444 which is better than the pruned tree model. Although decision trees are able to be utilized in both regression and classification and are easier to interpret compared to other machine learning models, models like Logistic Regression can provide a better predictive accuracy depending on the data set. I also noticed that the results were extremely close amongst most of the models so I

decided to run different seeds to get a better understanding of the average performance. However, I doubled checked with 3 different seeds (shown below) and it was interesting to see how the Logistic Regression outperformed in each instance and most likely more if additional seeds were run. This leads me to believe that the Logistic Regression model would provide consistent accurate results compared to the pruned decision trees but the decision trees are a lot easier to interpret for audiences. In the future, I plan on re-attempting this type of problem set for my own work and hopefully will be able to better understand each classification model's strengths in a personal work-based application.

Seed	666	999	777
Logistic Regression	0.151852	0.174074	0.174074
Tree model	0.166667	0.185185	0.2